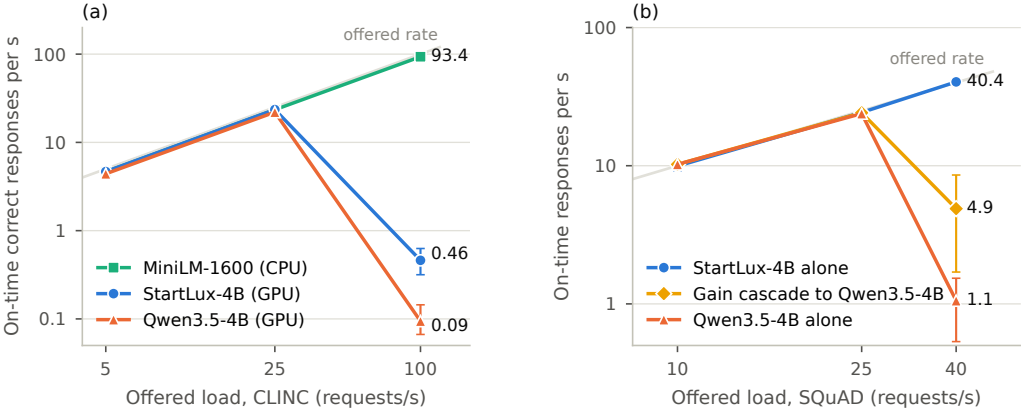
(a)
offered rate
93.4
0.46
0.09
MiniLM-1600 (CPU)
StartLux-4B (GPU)
Qwen3.5-4B (GPU)
On-time correct responses per s
100
10
1
0.1
5
25
100
Offered load, CLINC (requests/s)
(b)
offered rate
40.4
4.9
1.1
StartLux-4B alone
Gain cascade to Qwen3.5-4B
Qwen3.5-4B alone
On-time responses per s
100
10
1
10
25
40
Offered load, SQuAD (requests/s)